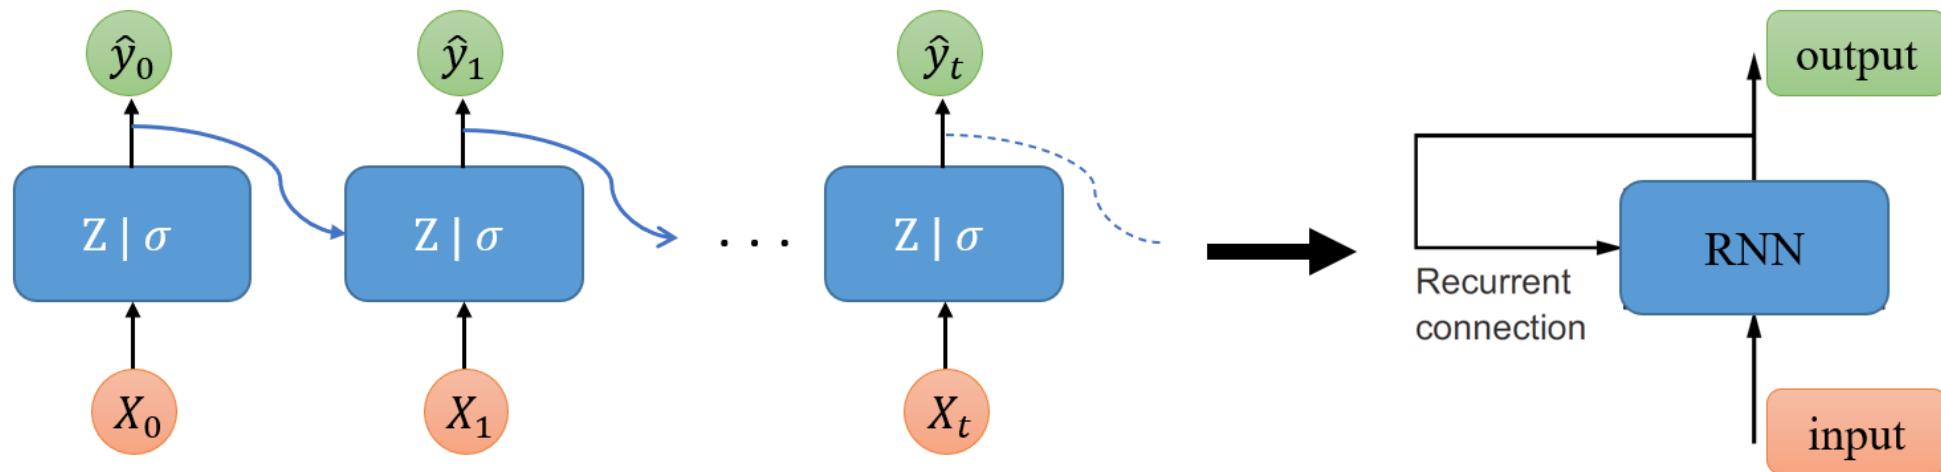


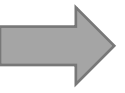
# Module 7 – Part I

## Deep Sequence Modeling

### Recurrent Neural Networks (RNN)

---

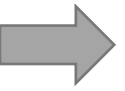




# Road map!

- Module 1- Introduction to Deep Forecasting
- Module 2- Setting up Deep Forecasting Environment
- Module 3- Exponential Smoothing
- Module 4- ARIMA models
- Module 5- Machine Learning for Time series Forecasting
- Module 6- Deep Neural Networks
- **Module 7- Deep Sequence Modeling (RNN, LSTM)**
- Module 8- Transformers (Attention is all you need!)
- Module 9- Prophet and Neural Prophet





# Different kinds of sequence data

Sequence data refers to any data that has a specific order or sequence to it!

- **Time series data:** sequence of data measured at regular intervals over time. (stock prices, weather patterns, medical records, ...)
- **Text data:** sequence of data that is composed of words, sentences, or paragraphs. (tweets, news articles, product reviews, ...)
- **Audio data:** sequence of data that is recorded or generated as sound waves. (speech recordings, music tracks, ...)
- **Video data:** sequence of data that is represented as a sequence of images or frames. (movie clips, surveillance footage, ...)





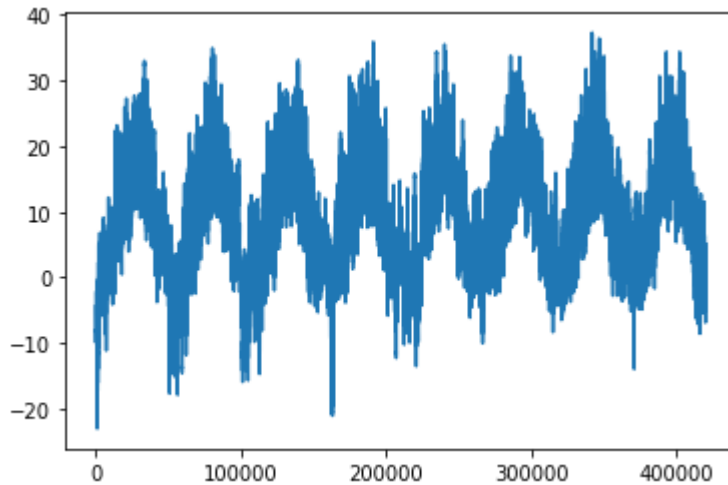
# Different kinds of timeseries tasks

---

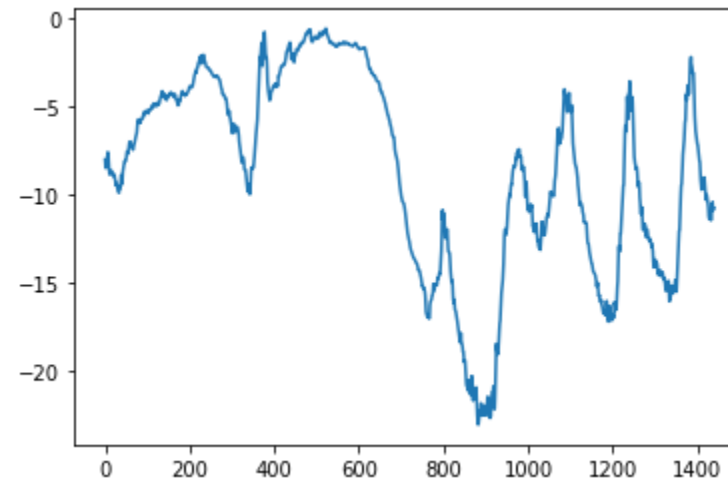
- **Forecasting**: Predicting what happens next  
Electricity consumption, temperature forecast, stock prediction, ...
- **Classification**: Assign one or more labels to a time series  
Classify website visitors as human or bot based on their activity
- **Event detection**: Identify the occurrence of a specific expected event within a continuous data stream  
Hotword detection: “Ok Google” or “Hey Siri”
- **Anomaly Detection**: Detect anything unusual happening within a continuous data stream  
Unusual activity on a network, ...

# ➔ A simple timeseries example

- A temperature forecasting example: [deep-learning-with-python-notebooks](#)
- Predicting the temperature 24 hours in the future
  - Target: temperature
  - Features: 14 different variables including pressure, humidity, wind direction and etc
  - Data recorded every 10 minutes from 2009-2016



Temperature between 2009-2016



Temperature in the first 10 days:  $10 \times 24 \times 6 = 1440$



# Preparing the data

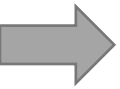
- Given the **previous 5 days (120 hours)** and samples **once per hour**, can we predict temperature **in 24 hours (after the end of the sequence)**?
- Data batches:
  - Sequence length = 120
  - $[1, 2, 3, \dots, 120][144]$
  - $[2, 3, 4, \dots, 121][145]$
  - $[3, 4, 5, \dots, 122][146]$
  - Batch size: 256 of these samples are shuffled and batched
  - Sample shape: (256, 120, 14)
  - Target shape: (256,)



# Naïve forecaster: common-sense baseline

- Temperature 24 hours from now = Temperature right now
- This is our random walk with no drift forecaster.
- Performance:
  - Validation MAE = 2.44 degrees Celsius
  - Test MAE = **2.62** degrees Celsius
  - The baseline model is off by about 2.5 degrees on average. Not bad!!





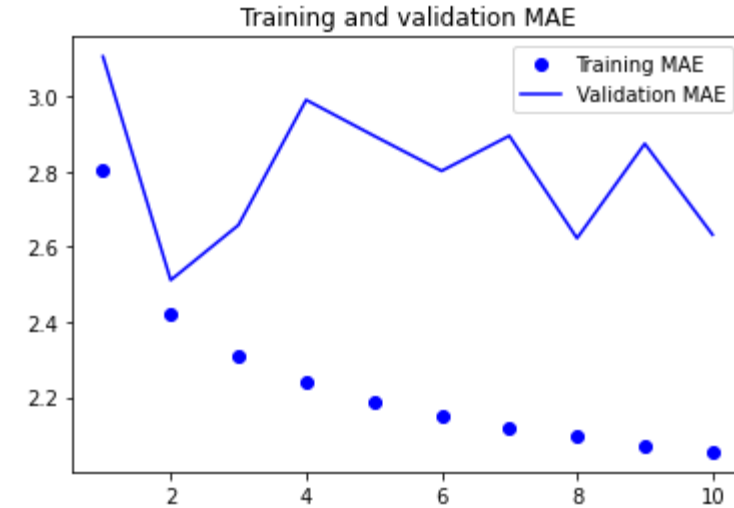
# Let's try DNN (Deep Neural Networks)

```
inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))
x = layers.Flatten()(inputs)
x = layers.Dense(16, activation="relu")(x)
outputs = layers.Dense(1)(x)
model = keras.Model(inputs, outputs)
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 120, 14)]	0
flatten (Flatten)	(None, 1680)	0
dense (Dense)	(None, 16)	26896
dense_1 (Dense)	(None, 1)	17

---

=====  
Total params: 26,913  
Trainable params: 26,913  
Non-trainable params: 0  
=====



- Test MAE = **2.62** degrees Celsius
- No improvement!!
- Flattening a timeseries data is not a good idea!



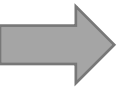


# Let's try CNN (Convolutional Neural Networks)

- **Motivation:** Maybe a temporal convnet could **reuse the same representations** across different days, much like a spatial convnet can reuse the same representations across different locations in an image!

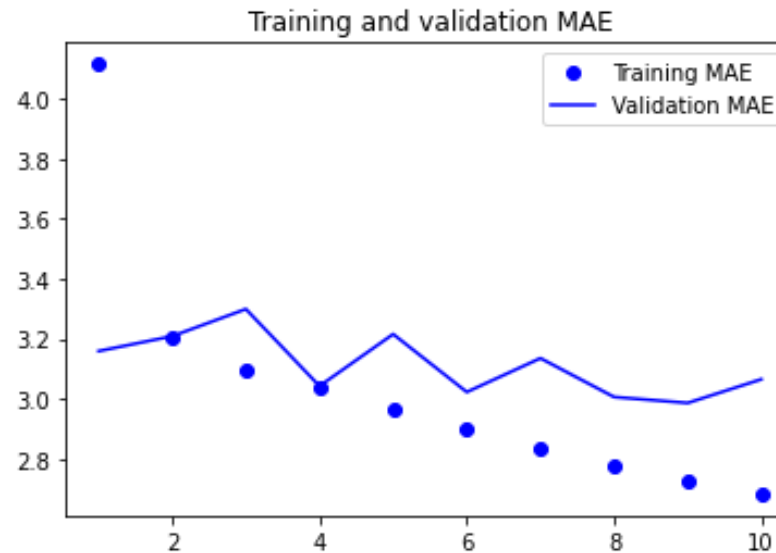
```
inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))
x = layers.Conv1D(8, 24, activation="relu")(inputs)
x = layers.MaxPooling1D(2)(x)
x = layers.Conv1D(8, 12, activation="relu")(x)
x = layers.MaxPooling1D(2)(x)
x = layers.Conv1D(8, 6, activation="relu")(x)
x = layers.GlobalAveragePooling1D()(x)
outputs = layers.Dense(1)(x)
model = keras.Model(inputs, outputs)
```

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 120, 14)]	0
conv1d (Conv1D)	(None, 97, 8)	2696
max_pooling1d (MaxPooling1D)	(None, 48, 8)	0
conv1d_1 (Conv1D)	(None, 37, 8)	776
max_pooling1d_1 (MaxPooling1D)	(None, 18, 8)	0
conv1d_2 (Conv1D)	(None, 13, 8)	392
global_average_pooling1d (GlobalAveragePooling1D)	(None, 8)	0
dense_2 (Dense)	(None, 1)	9
Total params: 3,873		
Trainable params: 3,873		
Non-trainable params: 0		



# CNN performance

- Test MAE = **3.10** degrees Celsius
- Even worse than the densely connected model!!!
  - CNN treats every segment of the data the same way!
  - Pooling layers are destroying order information.



# → Sequence Modeling

---

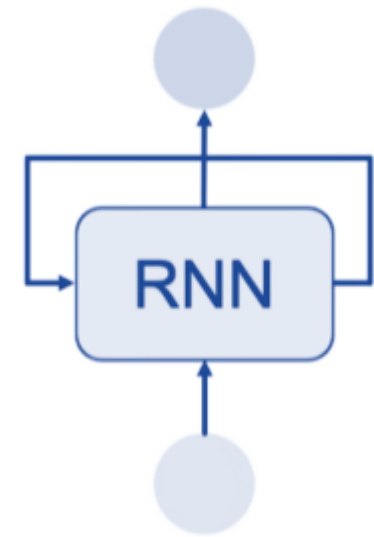
To model sequence data efficiently, we need a new architecture that:

- Preserve the **order**
- Account for **long-term dependencies**
- Handle **input-length**
- **Share parameters** across the sequence

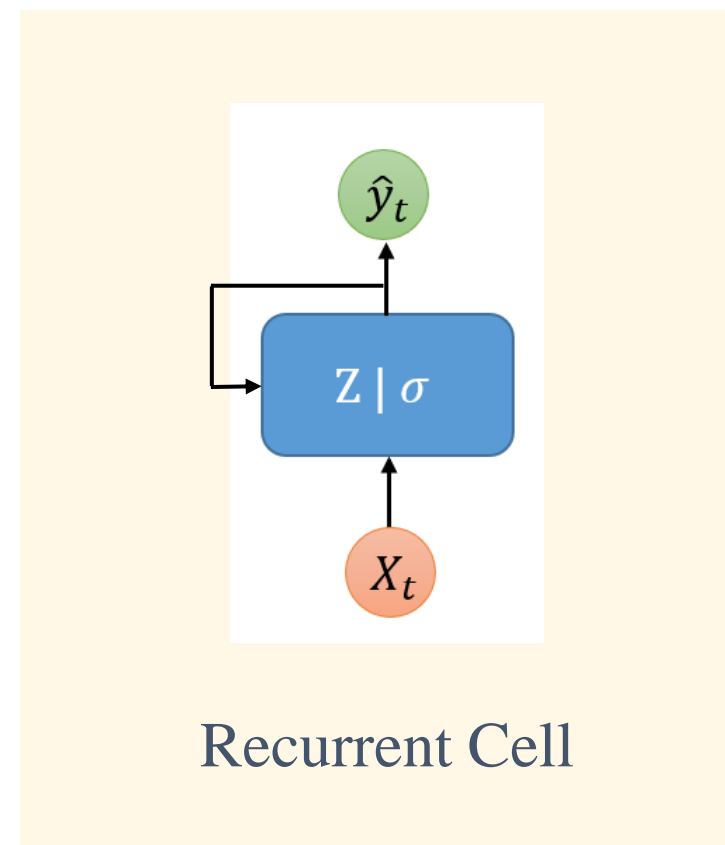
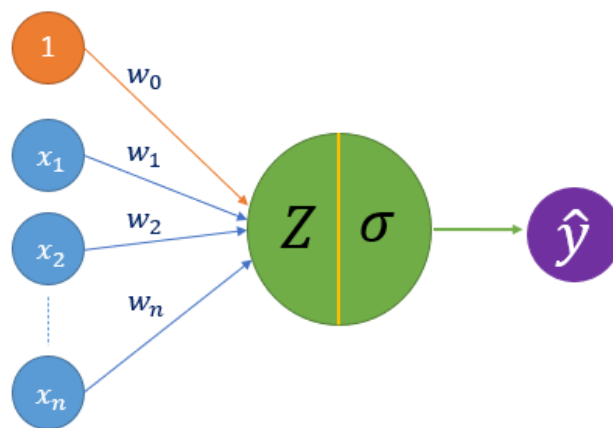
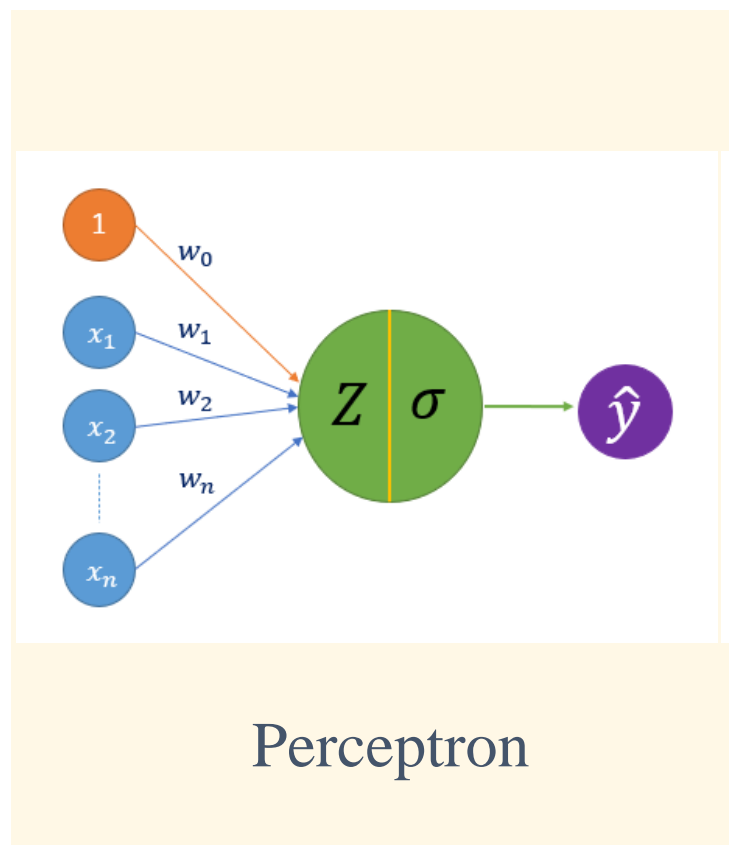


# What is RNN (Recurrent Neural Network)?

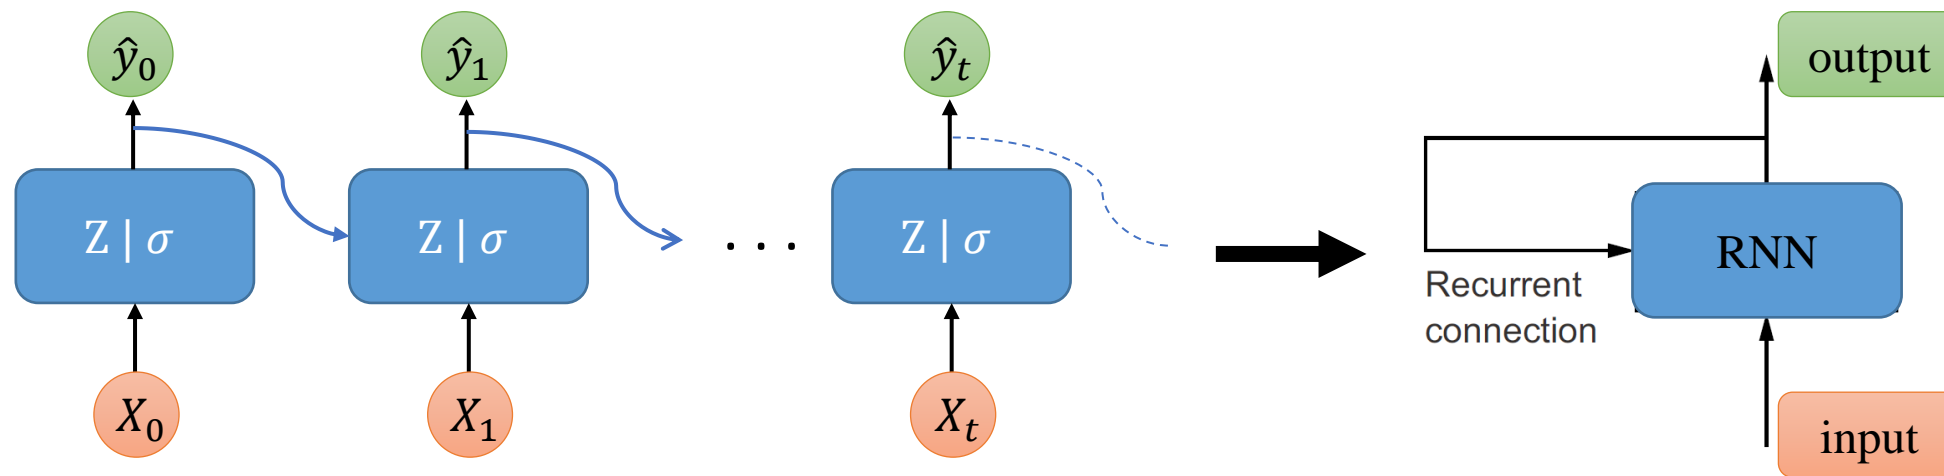
- The architecture of RNNs is inspired by the way **biological intelligence** processes information **incrementally** while **maintaining an internal model** of what it is processing.
- This ability to **remember previous inputs** and incorporate them into the current output allows RNNs to model sequential data.
- RNN maintains a **state** that contains information relative to what it has seen so far
- RNNs can be thought of as neural networks with an **internal loop**, which allows them to process sequences of varying lengths and learn from temporal dependencies.



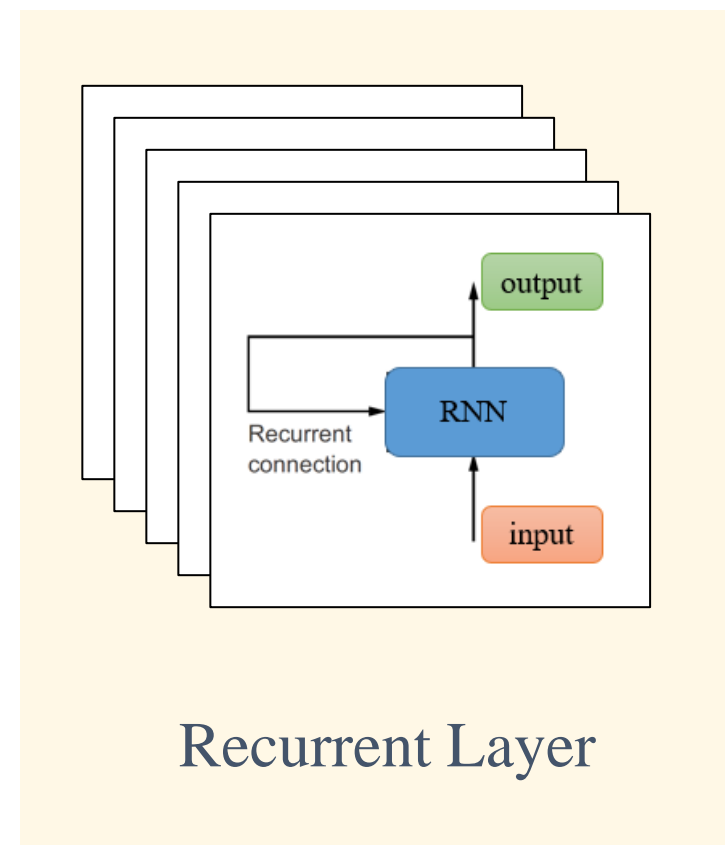
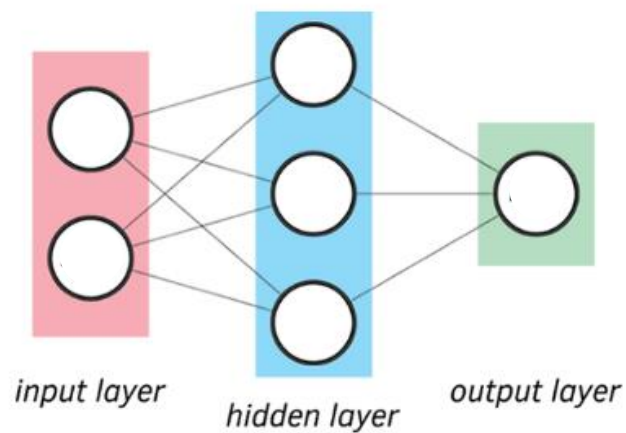
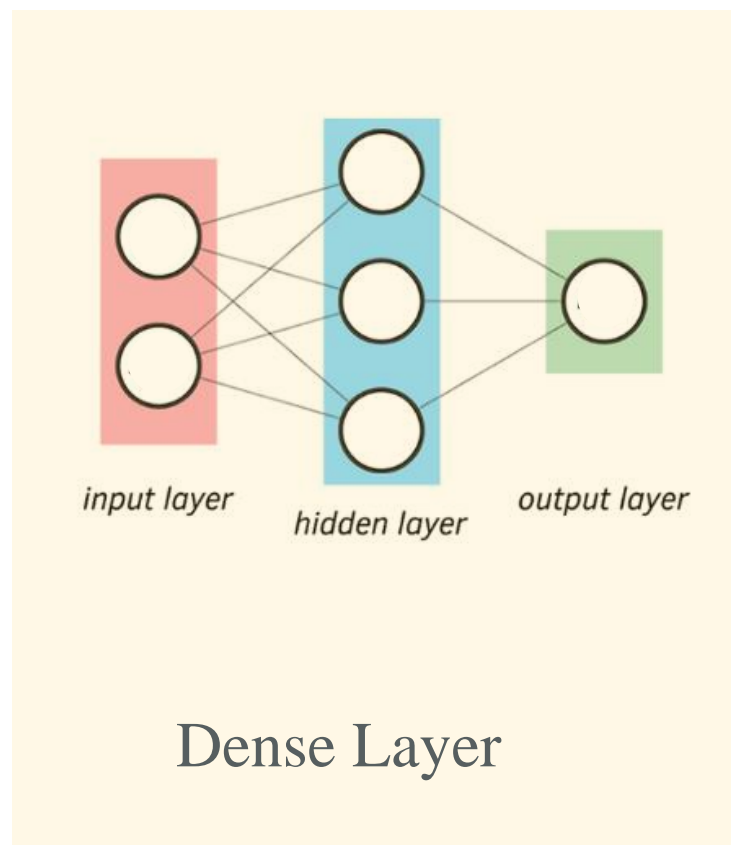
# ➔ Perceptron vs Recurrent Cell



# → Unrolling the Recurrent Cell

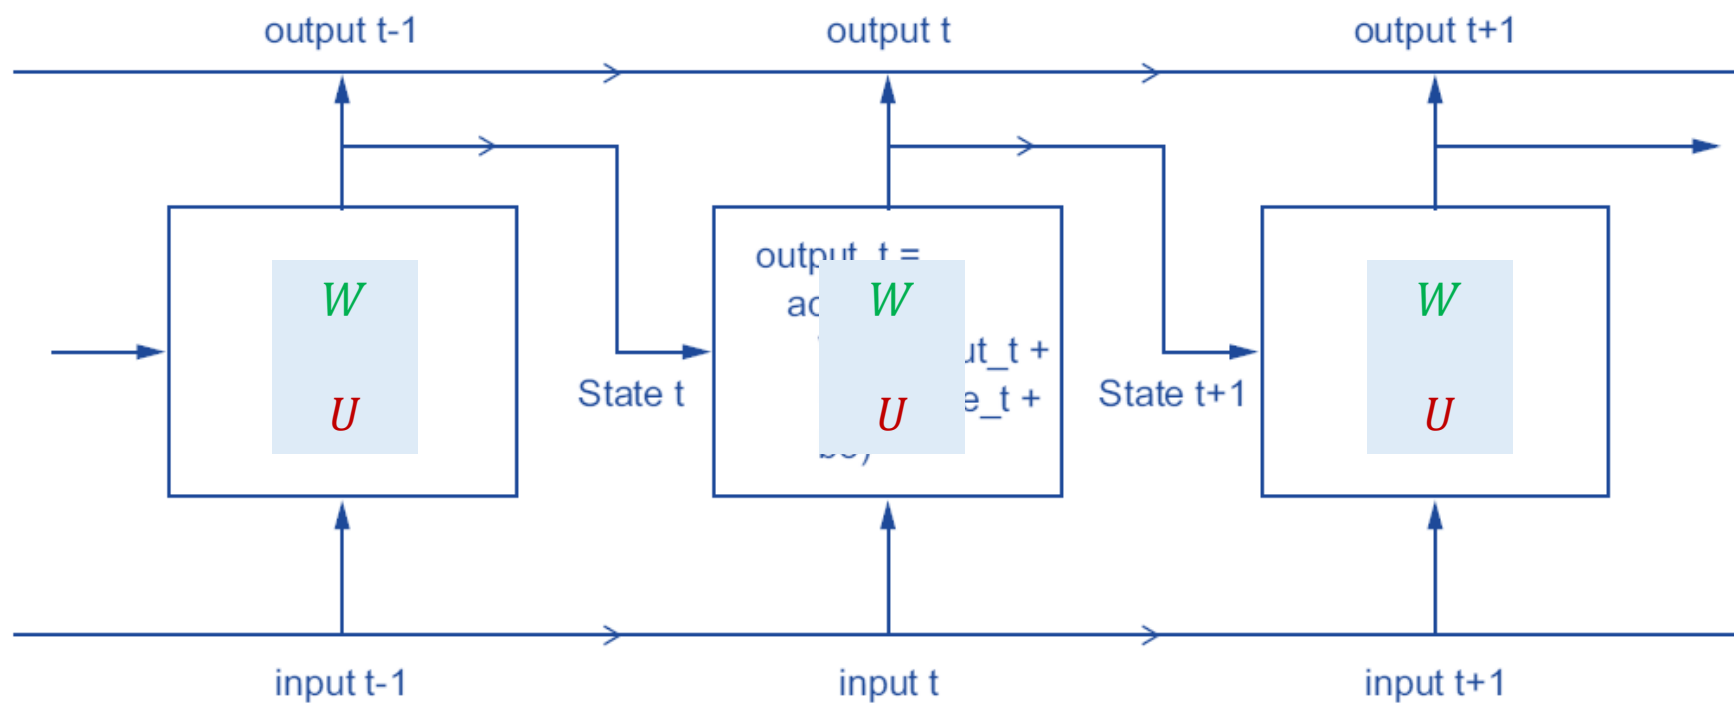


# ➔ Dense Layer vs Recurrent Layer



# → Inside the Recurrent Cell

$$output_t = f(input_t, State_t)$$

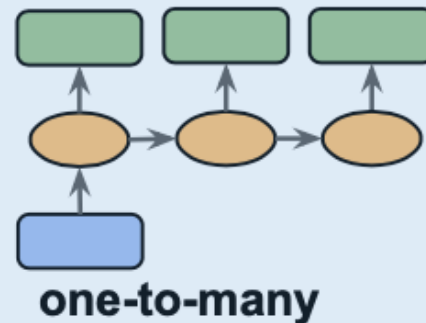
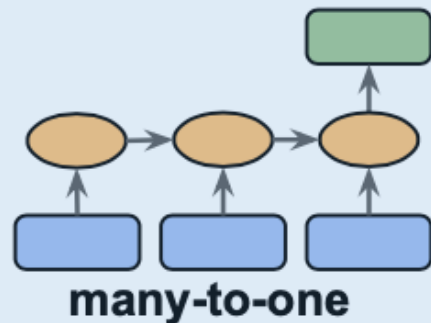
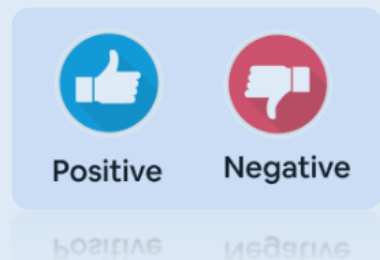


$$s_{t+1} = activation(WX_t + Us_t + b)$$

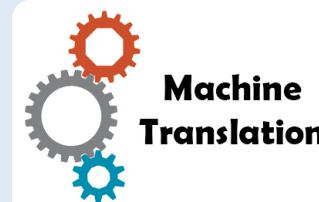
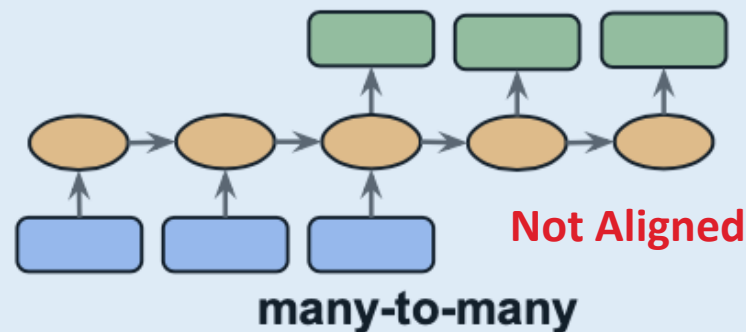
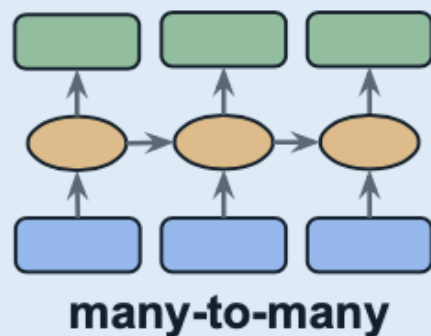




# RNN architectures



a train traveling down a track next to a forest.





# How does RNN learn representations?

- Backpropagation Through Time (BPTT)

- $\frac{\partial J}{\partial P}$  P are the parameters

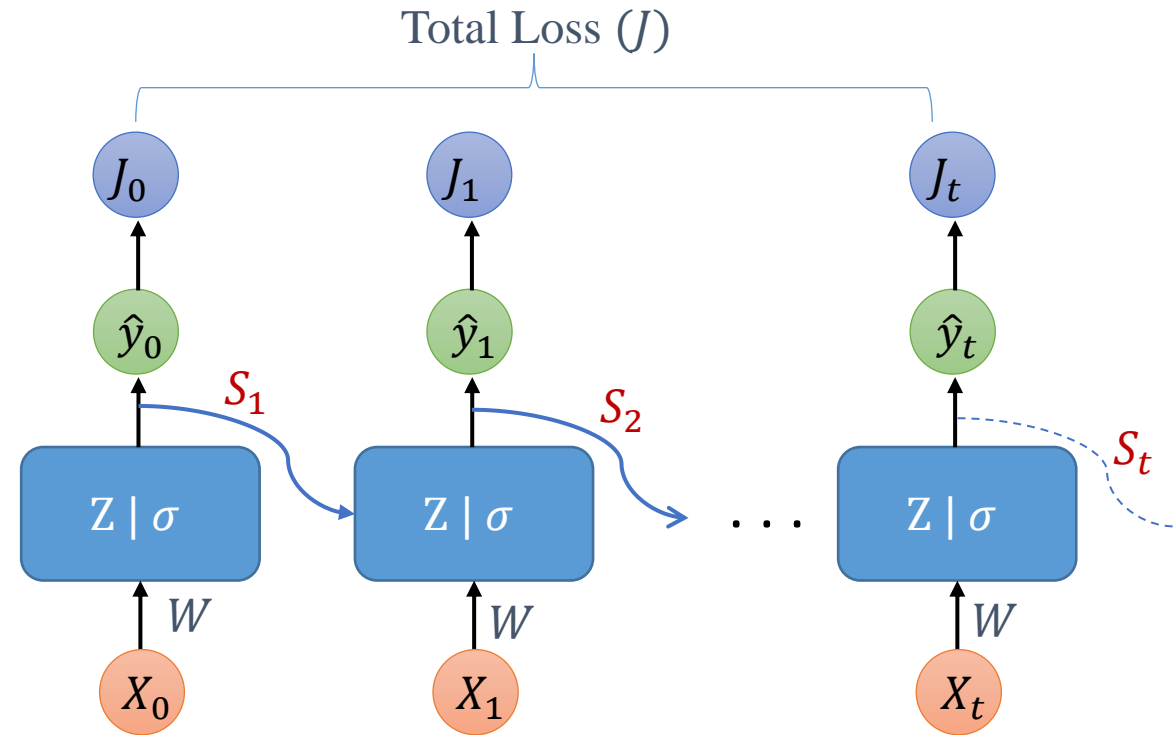
- $\frac{\partial J}{\partial W} = \frac{\partial J_0}{\partial W} + \frac{\partial J_1}{\partial W} + \dots$

- $\frac{\partial J_0}{\partial W} = \frac{\partial J_0}{\partial y_0} \frac{\partial y_0}{\partial s_0} \frac{\partial s_0}{\partial W}$

- $\frac{\partial J_1}{\partial W} = \frac{\partial J_1}{\partial y_1} \frac{\partial y_1}{\partial s_1} \frac{\partial s_1}{\partial W}$  ,  $\frac{\partial s_1}{\partial W} = \frac{\partial s_1}{\partial s_0} \frac{\partial s_0}{\partial W}$

- ...

- $\frac{\partial J_t}{\partial W} = \sum_{k=0}^t \frac{\partial J_t}{\partial y_t} \frac{\partial y_t}{\partial s_t} \frac{\partial s_t}{\partial s_k} \frac{\partial s_k}{\partial W}$



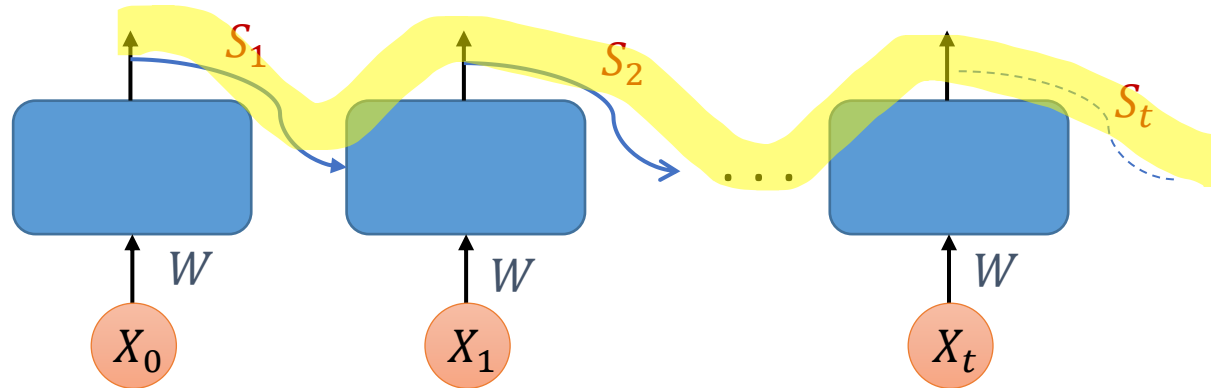
# → Vanishing Gradient Problem

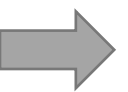
- As the time horizon gets bigger, this product gets longer and longer.
- We are multiplying a lot of small numbers → smaller gradients → biased parameters unable to capture long term dependencies.

$$\bullet \frac{\partial J_t}{\partial W} = \sum_{k=0}^t \frac{\partial J_t}{\partial y_t} \frac{\partial y_t}{\partial S_t} \frac{\partial S_t}{\partial S_k} \frac{\partial S_k}{\partial W}$$

$$\bullet \frac{\partial S_{10}}{\partial S_0} = \frac{\partial S_{10}}{\partial S_9} \frac{\partial S_9}{\partial S_8} \frac{\partial S_8}{\partial S_7} \frac{\partial S_7}{\partial S_6} \cdots \frac{\partial S_1}{\partial S_0}$$

$$S_t = \text{activation}(W X_{t-1} + U S_{t-1})$$





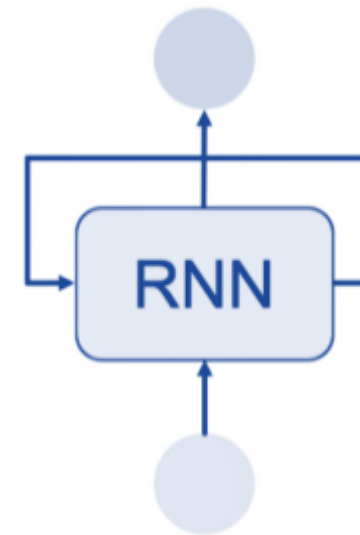
# Beyond RNN

RNN can handle the following sequence modeling criteria:

- Preserve the **order**
- Handle **input-length**
- **Share parameters** across the sequence

RNN limitations:

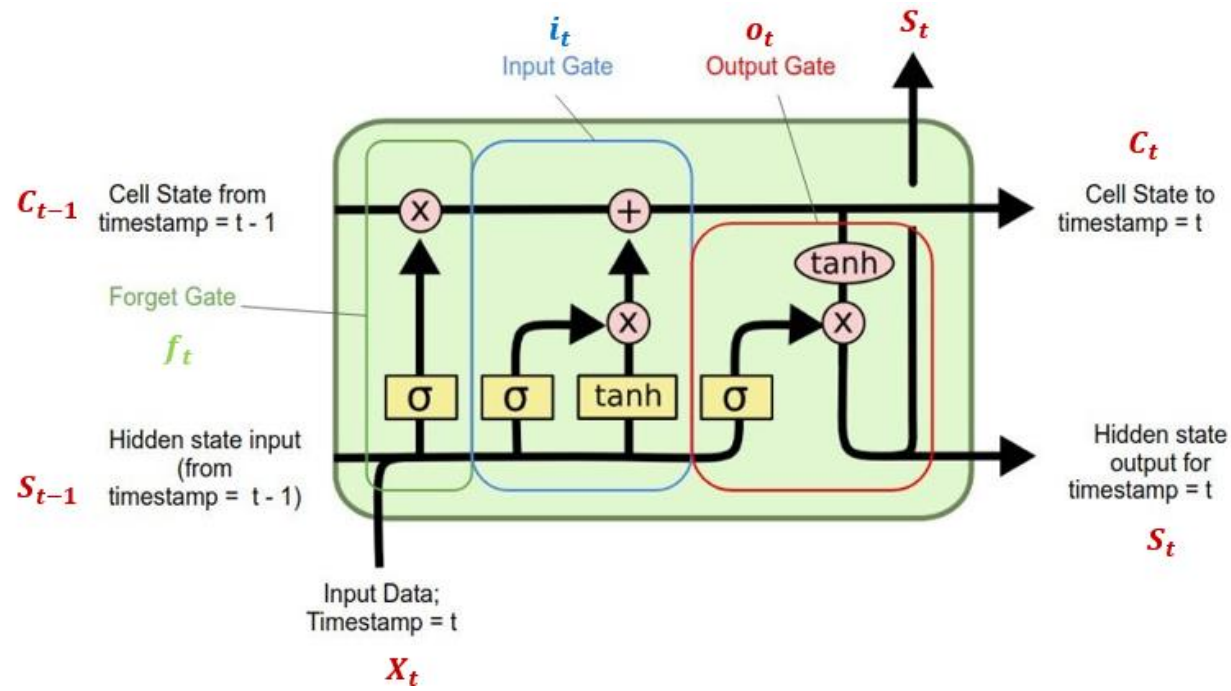
- Does not account for **long-term dependencies** (only remember short term history )
- Vanishing Gradient Problem

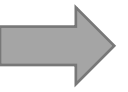


# Module 7 – Part II

## Deep Sequence Modeling (Gated cells, LSTM)

---

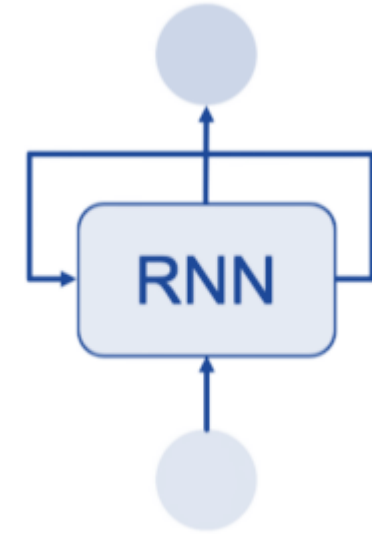




# Beyond RNN

RNN can handle the following sequence modeling criteria:

- Preserve the **order**
- Handle **input-length**
- **Share parameters** across the sequence



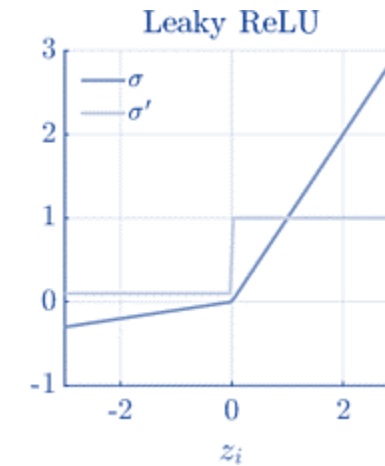
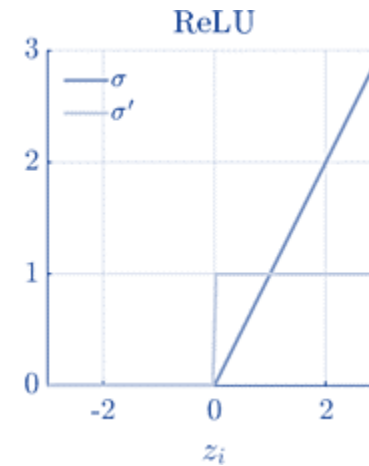
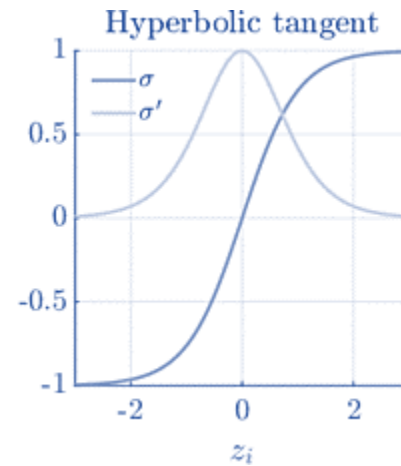
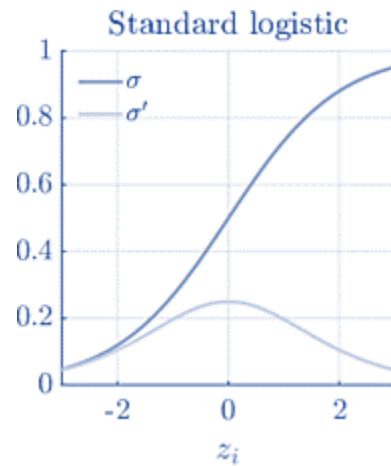
RNN limitations:

- Does not account for **long-term dependencies** (only remember short term history )
- Vanishing Gradient Problem



# How to solve vanishing gradient problem

1. Use **Activation Function** that prevents fast shrinkage of gradient



$$S_t = \text{activation}(\textcolor{green}{W}X_{t-1} + \textcolor{red}{U}s_{t-1})$$



# How to solve vanishing gradient problem

---

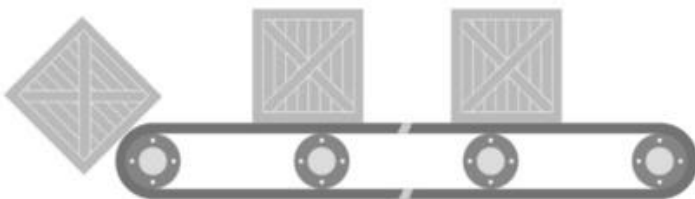
1. Use **Activation Function** that prevents fast shrinkage of gradient
2. Use **weight initialization** techniques that ensure that the initial weights are not too small
3. Use **gradient clipping** which limits the magnitude of the gradients from becoming too small (vanishing gradient) or too large (exploding gradient)
4. Use **batch normalization**, which normalizes the input to each layer and helps to reduce the range of activation values and thus the likelihood of vanishing gradients.
5. Use a different **optimization algorithm** that is more resilient to vanishing gradients, such as Adam or RMSprop.
6. **Gated cells:** Use some sort of **skip connections**, which allow gradients to bypass some of the layers in the network and thus prevent them from becoming too small.





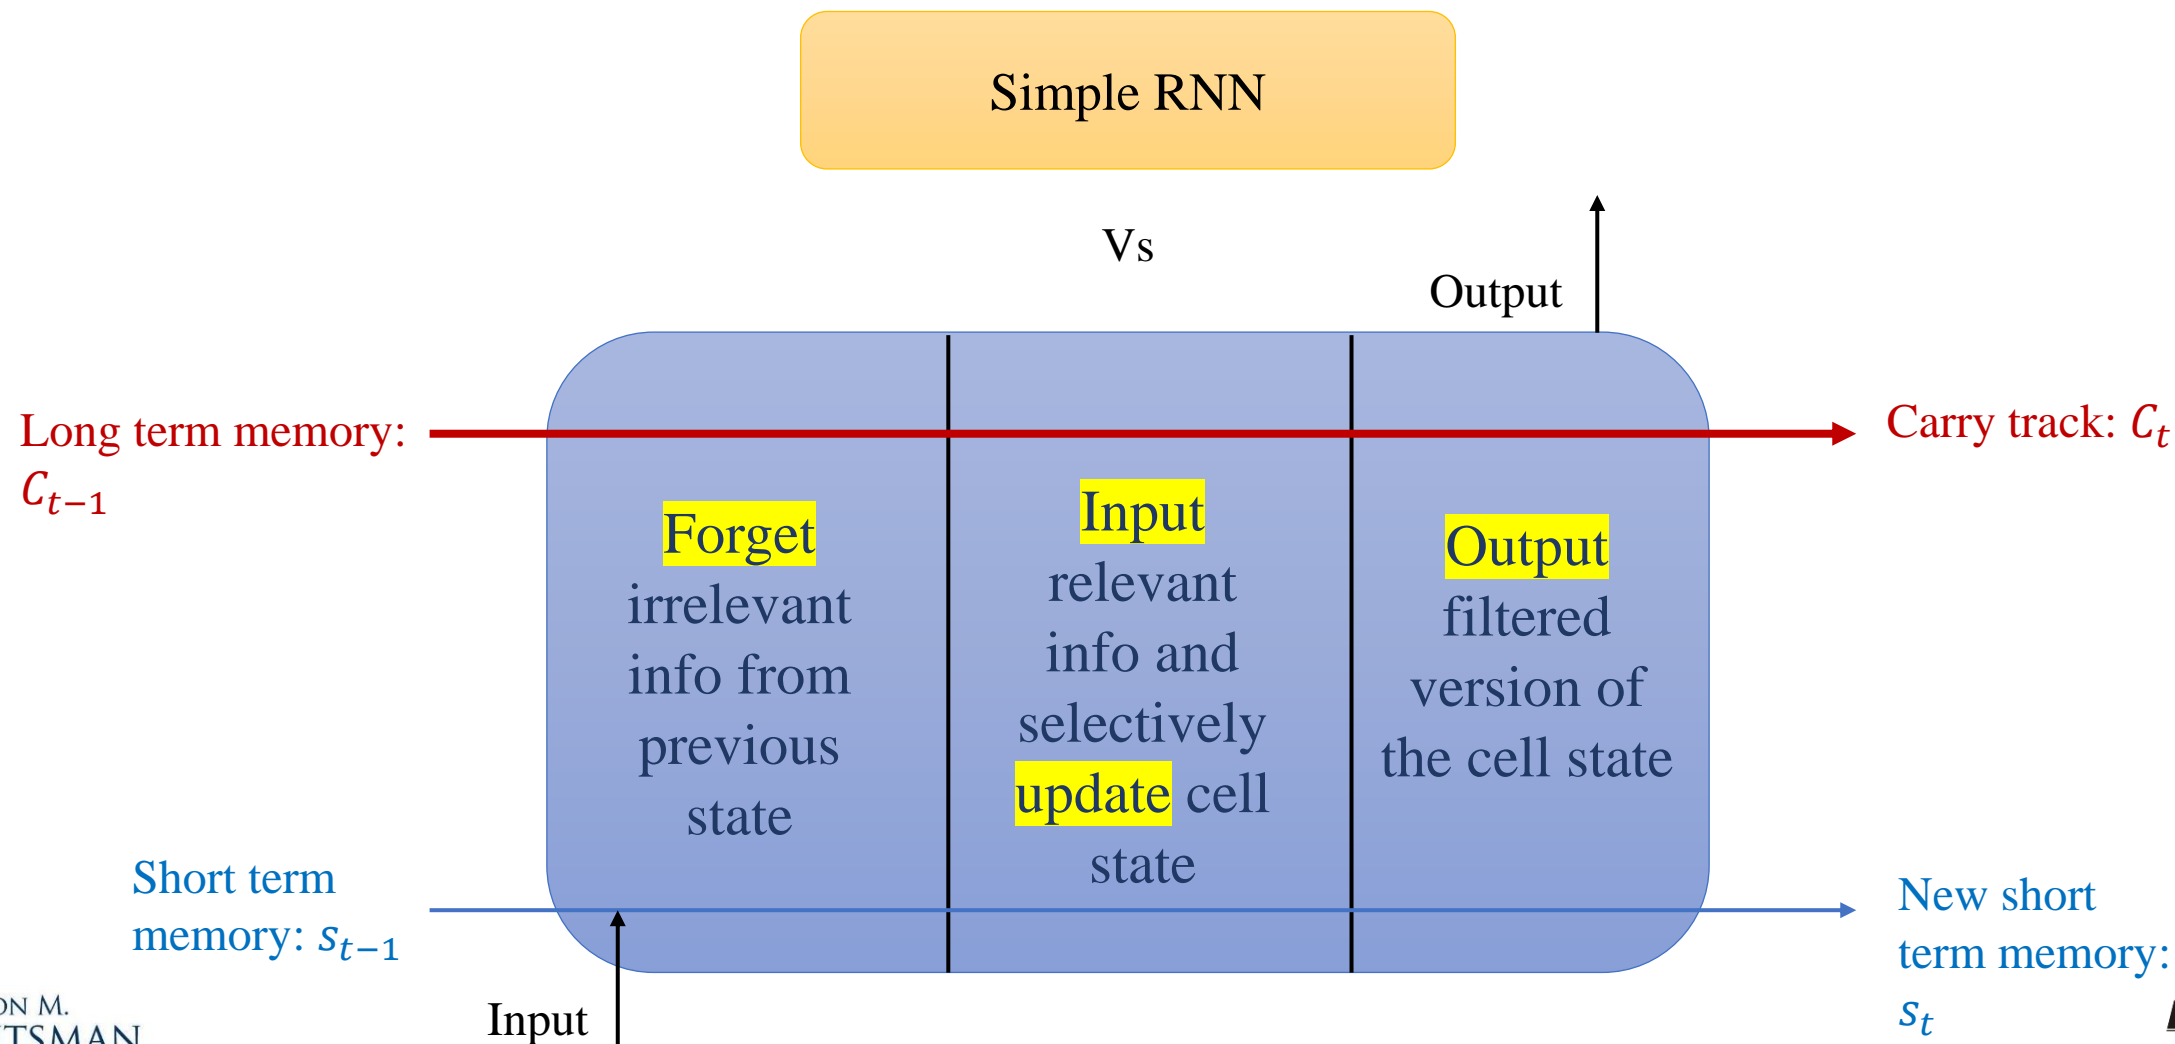
# Gated cells

- Instead of using a simple RNN cell, let's use a **more complex cell** with gates which **control the flow of information**.
- Think of a **conveyor belt** running parallel to the sequence being processed:
  - Information can jump on → transported to a later timestep → jump off when needed.
  - This is what a gated cell does! Analogous to **residual connections** we saw before.



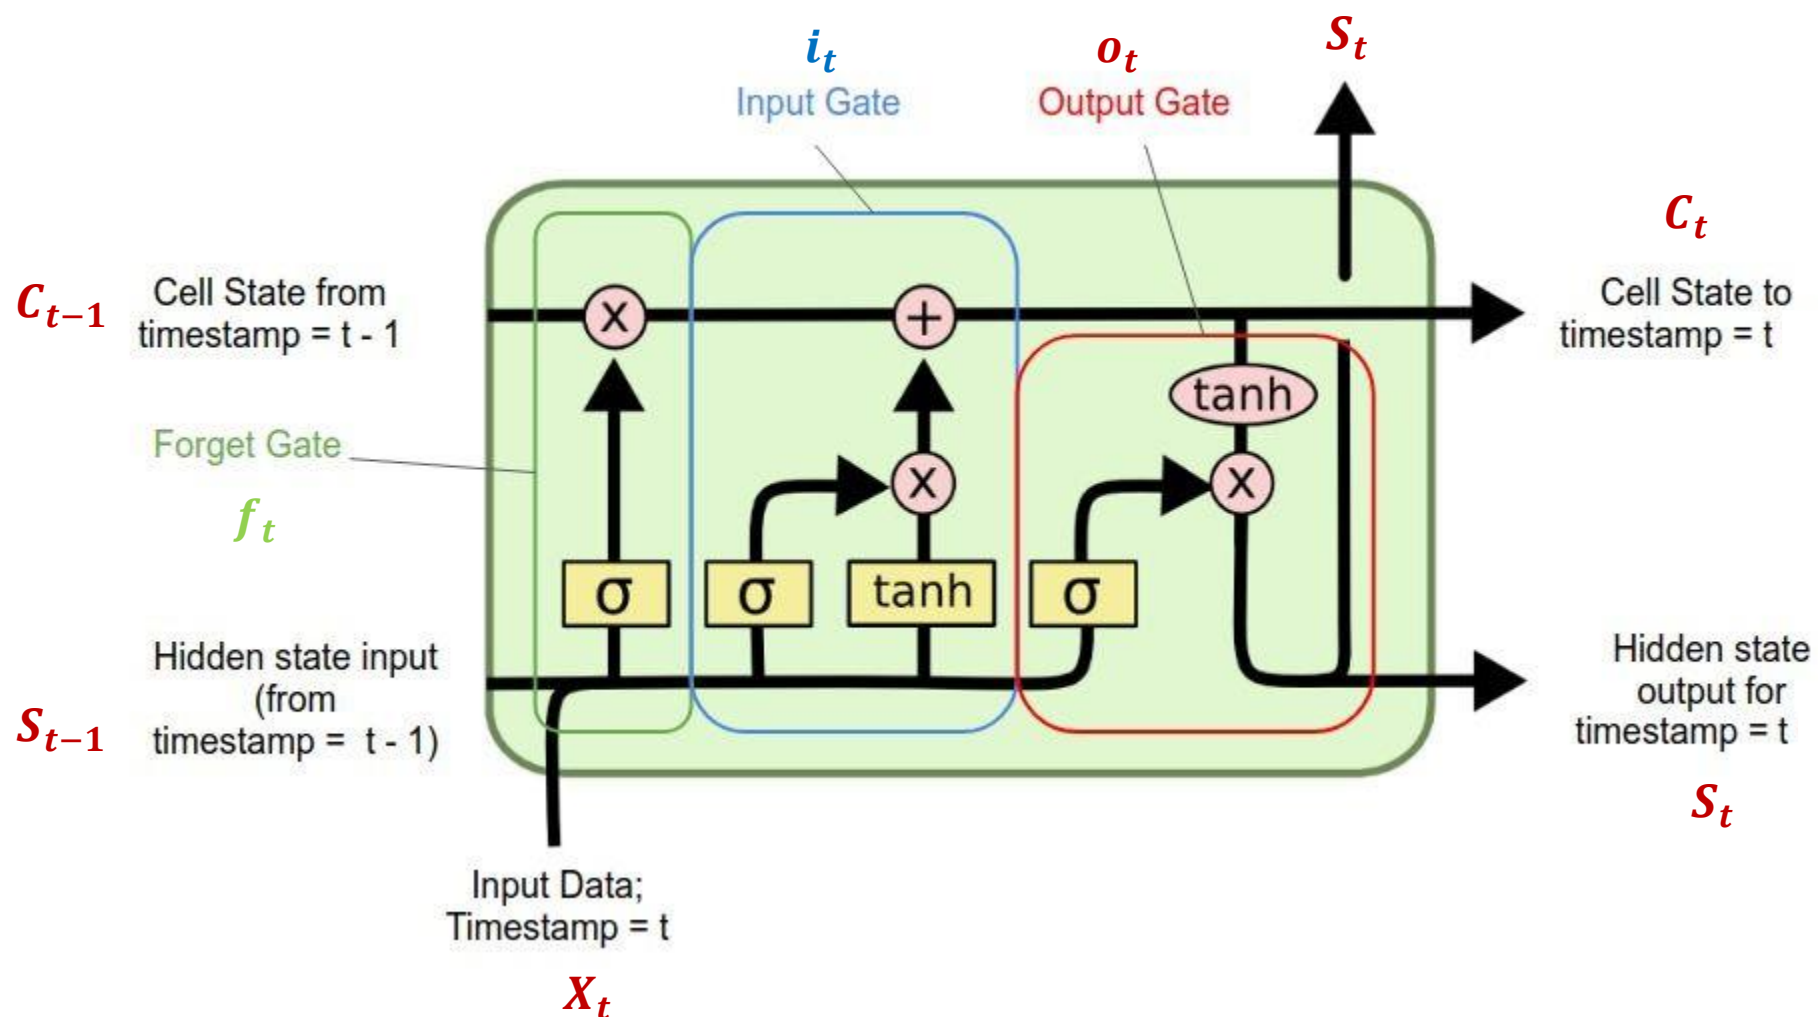
- Long Short-Term Memory (LSTM) and Gated Recurrent Unit (GRU) are two examples of gated cells that can **keep track of information throughout many timesteps**.

# → Inside the LSTM cell





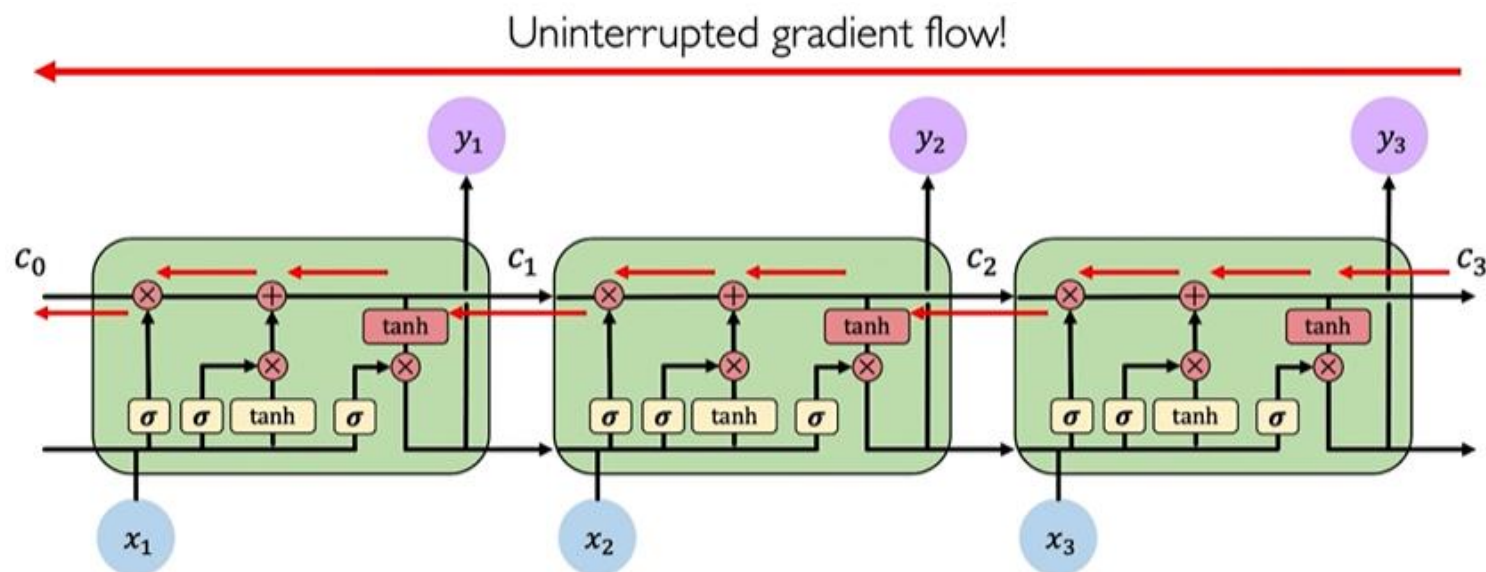
# LSTM details

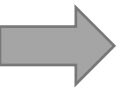




# LSTM takeaway

- LSTM uses gates to **regulate the information flow** (allows past information to be reinjected later)
- This new cell state (carry) can better capture **longer term dependencies**
- LSTM **fights** the vanishing gradient problem



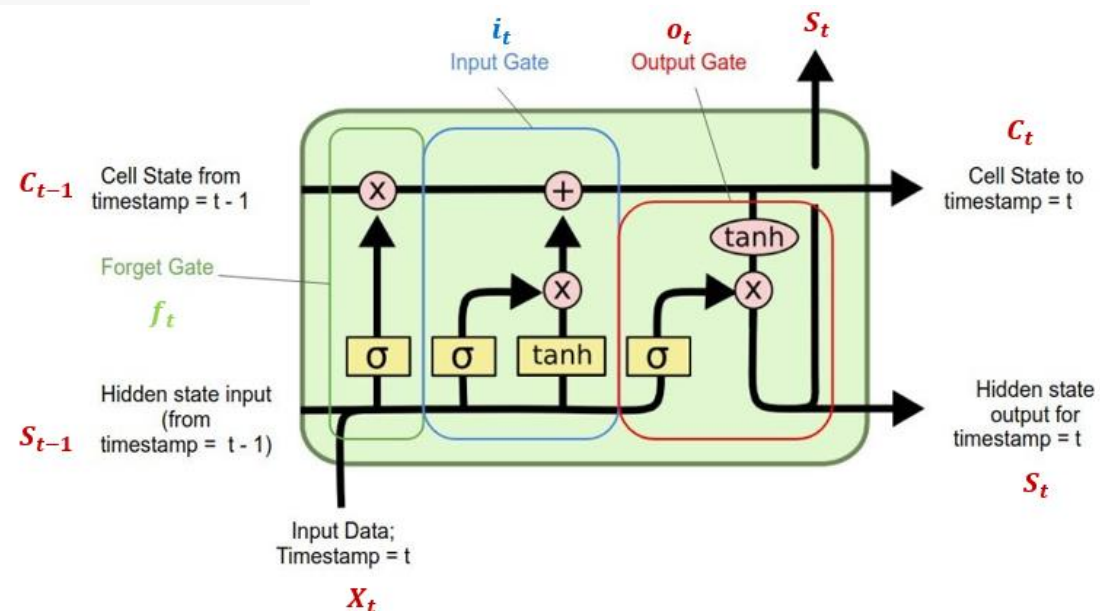


# Let's try LSTM on the temperature example

```
inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))
x = layers.LSTM(16)(inputs)
outputs = layers.Dense(1)(x)
model = keras.Model(inputs, outputs)
```

Layer (type)	Output Shape	Param #
input_3 (InputLayer)	[(None, 120, 14)]	0
lstm (LSTM)	(None, 16)	1984
dense_3 (Dense)	(None, 1)	17

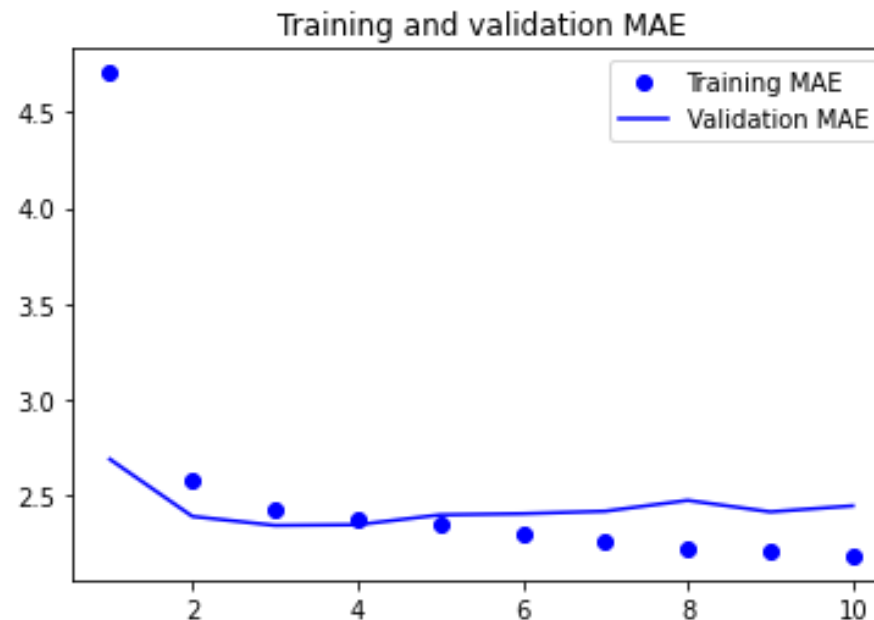
=====  
Total params: 2,001  
Trainable params: 2,001  
Non-trainable params: 0  
=====



```
output_t = activation(dot(state_t, Uo) + dot(input_t, Wo) + dot(c_t, Vo) + bo)
i_t = activation(dot(state_t, Ui) + dot(input_t, Wi) + bi)
f_t = activation(dot(state_t, Uf) + dot(input_t, Wf) + bf)
k_t = activation(dot(state_t, Uk) + dot(input_t, Wk) + bk)
```

# ➔ LSTM performance

- Baseline Test MAE = 2.62
- Simple LSTM Test MAE = 2.53
- Finally beat the naïve forecaster.
- Overfitting?



Can we do better?



# Improving the simple LSTM model

- We can improve the performance of the simple LSTM model by:
  1. **Recurrent Dropout** : use drop out to fight overfitting in the recurrent layers (in addition to drop out for the dense layers)
  2. **Stacking recurrent layers**: increase model complexity to boost representation power
  3. **Using bidirectional RNN**: processing the same information differently! Mostly used in NLP.



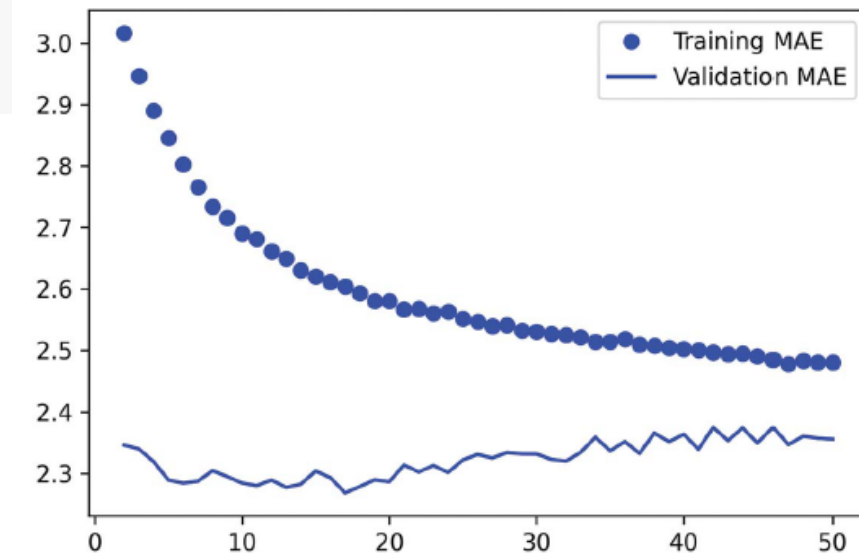


# → Recurrent Drop out

- The **same dropout pattern** should be applied at every timestep

```
inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))
x = layers.LSTM(32, recurrent_dropout=0.25)(inputs)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1)(x)
model = keras.Model(inputs, outputs)
```

- Baseline Test MAE = 2.62
- Simple LSTM, Test MAE = 2.53
- LSTM with dropout, Test MAE = 2.45

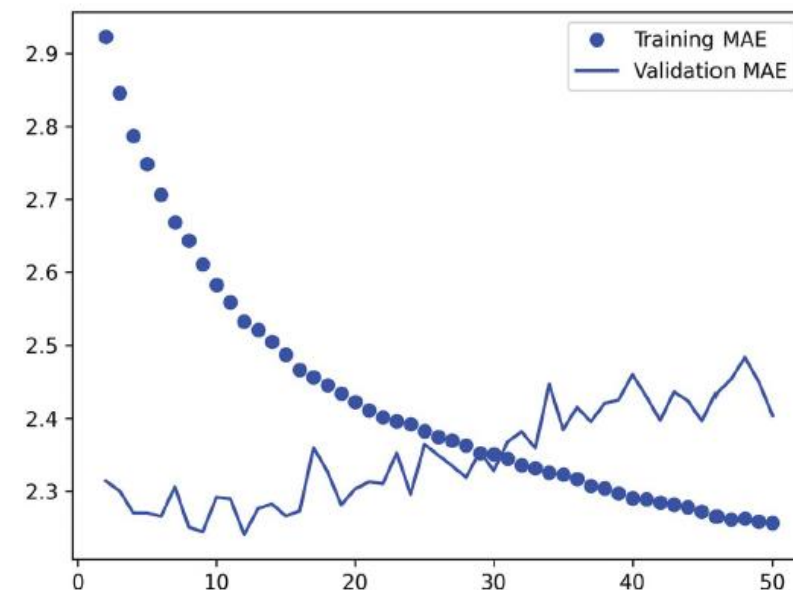


# → Stacking Recurrent Layers

- Let's train a dropout-regulated, stacked GRU model.
- GRU is a slightly simpler version (hence, faster) of LSTM architecture

```
inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))
x = layers.GRU(32, recurrent_dropout=0.5, return_sequences=True)(inputs)
x = layers.GRU(32, recurrent_dropout=0.5)(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1)(x)
model = keras.Model(inputs, outputs)
```

- Baseline Test MAE = 2.62
- Simple LSTM, Test MAE = 2.53
- Stacking GRU, Test MAE = 2.39

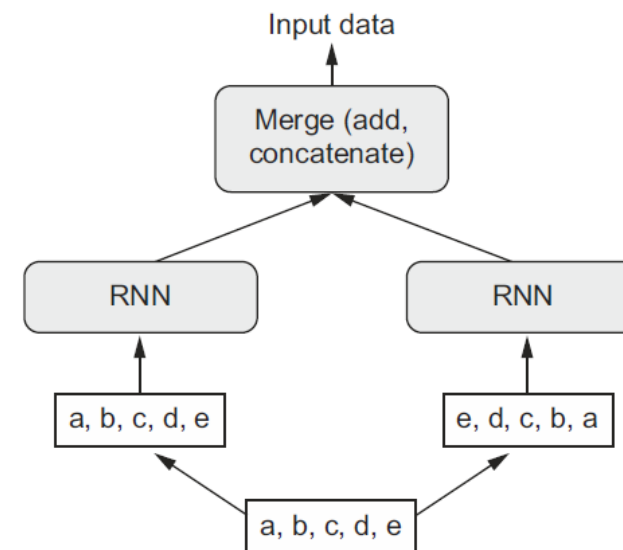


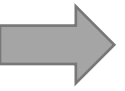


# Bidirectional RNN

- Bidirectional RNN process the input sequence both chronologically and antichronologically.
- Idea: capturing patterns (representations) that might be overlooked by a unidirectional RNN.
- For the temperature example, the bidirectional LSTM strongly underperforms even the common-sense baseline.

```
inputs = keras.Input(shape=(sequence_length, raw_data.shape[-1]))
x = layers.Bidirectional(layers.LSTM(16))(inputs)
outputs = layers.Dense(1)(x)
model = keras.Model(inputs, outputs)
```





# Final message

- Deep learning is more an art than science! Too many moving part!
  - Number of units in each recurrent layer
  - Number of stacked layers
  - Amount of dropout and recurrent dropout
  - Number of dense layers
  - Sequence horizon!
  - Optimizers, learning rates and etc
  - ....
- Apply RNN to datasets that past is a good predictor of the future! **Not the stock market!**

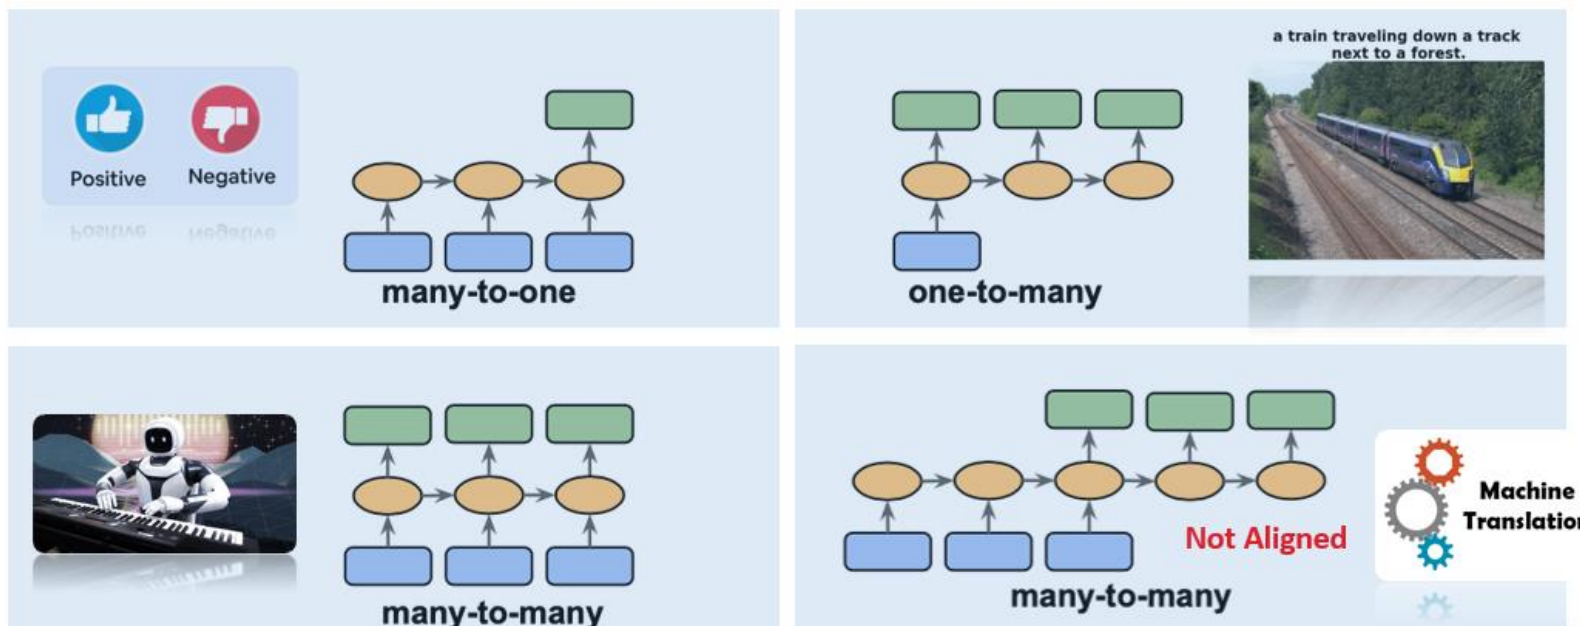


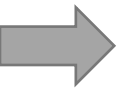
# Module 7 – Part III

## Deep Sequence Modeling for Text

### Natural Language Processing (NLP)

---





# Different kinds of sequence data

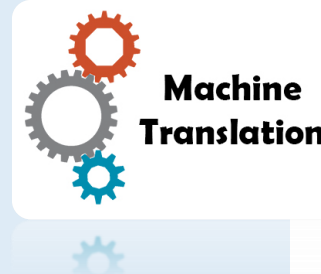
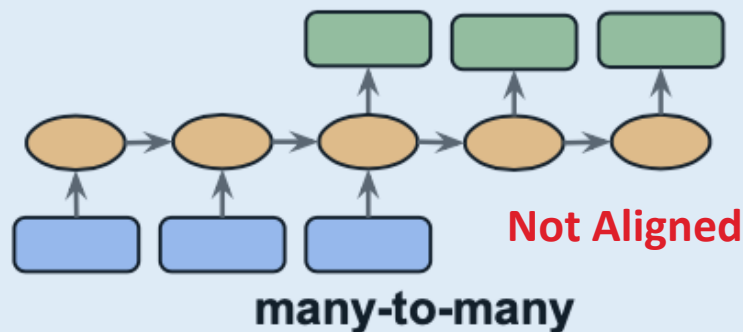
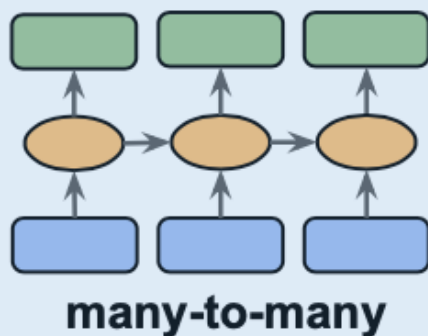
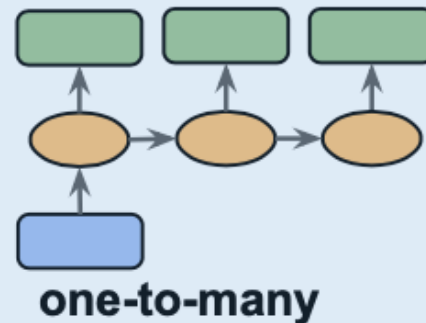
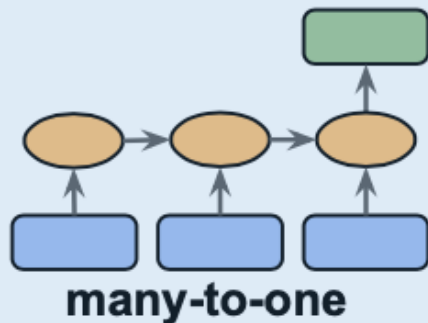
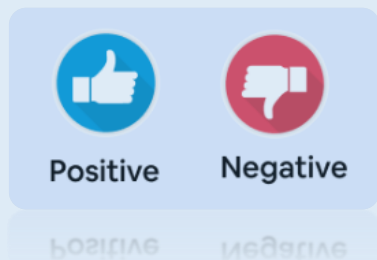
---

Sequence data refers to any data that has a specific order or sequence to it!

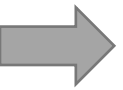
- **Time series data:** Time series data is a sequence of data points that are measured at regular intervals over time. (stock prices, weather patterns, medical records, ...)
- **Text data:** Text data refers to any type of data that is composed of words, sentences, or paragraphs. (tweets, news articles, product reviews, ...)
- **Audio data:** Audio data refers to any type of data that is recorded or generated as sound waves. (speech recordings, music tracks, ...)
- **Video data:** Video data refers to any type of data that is represented as a sequence of images or frames. (movie clips, surveillance footage, ...)



# RNN architectures







# Human language vs machine language

- In compute science: **Human** language → **Natural** language (shaped by an evolution process)
- **Machine** language designed for machines (XML, Assembly, ...)



Factors	Human Language	Machine Language
<b>Complexity</b>	Highly complex	Structured and logical
<b>Creativity</b>	Can be highly creative	Limited creativity
<b>Ambiguity</b>	High levels of ambiguity	Little to no ambiguity
<b>Learning</b>	Acquired naturally	Taught through programming
<b>Adaptability</b>	Highly adaptable	Fixed and rigid



# ➔ Natural Language Processing

- Making sense of human language through algorithms is a big deal!
- NLP is a subfield of computer science and AI that deals with the interaction between computers and human language
- Traditional NLP started with finding **handcrafted rules** to **understand** human language (1960-1990)
- Modern NLP is all about using machine learning to automate the search for these rules and enabling computers to do the followings:
  - Sentiment analysis
  - Machine translation
  - Content filtering
  - Text classification and ...

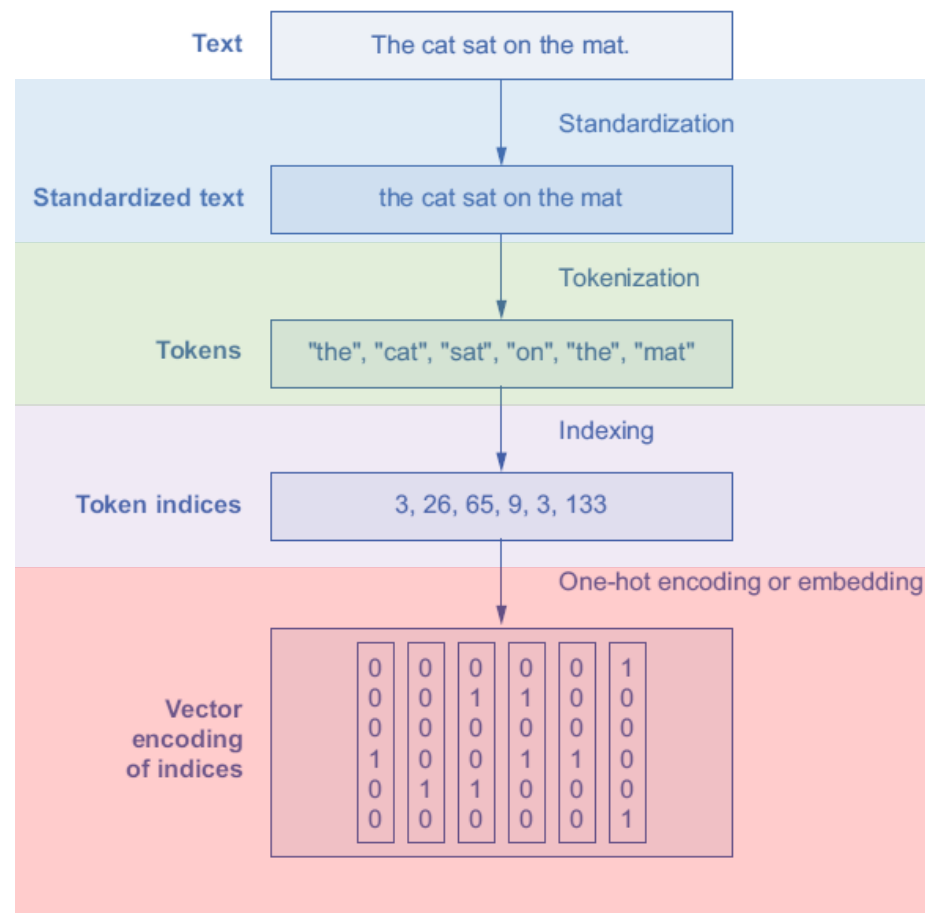




# Preparing text data

- **Text Vectorization**: the process of transforming text into numeric values.

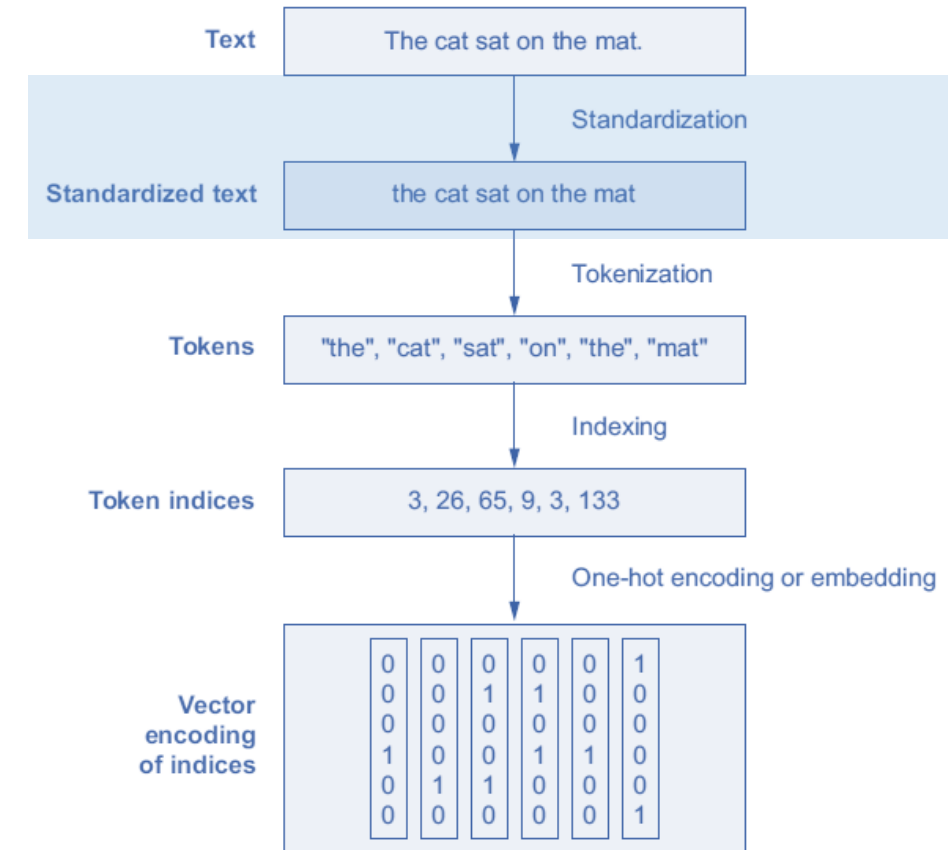
1. Standardization
2. Tokenization
3. Indexing
4. Encoding





# Text Vectorization: **standardization**, tokenization, indexing, encoding

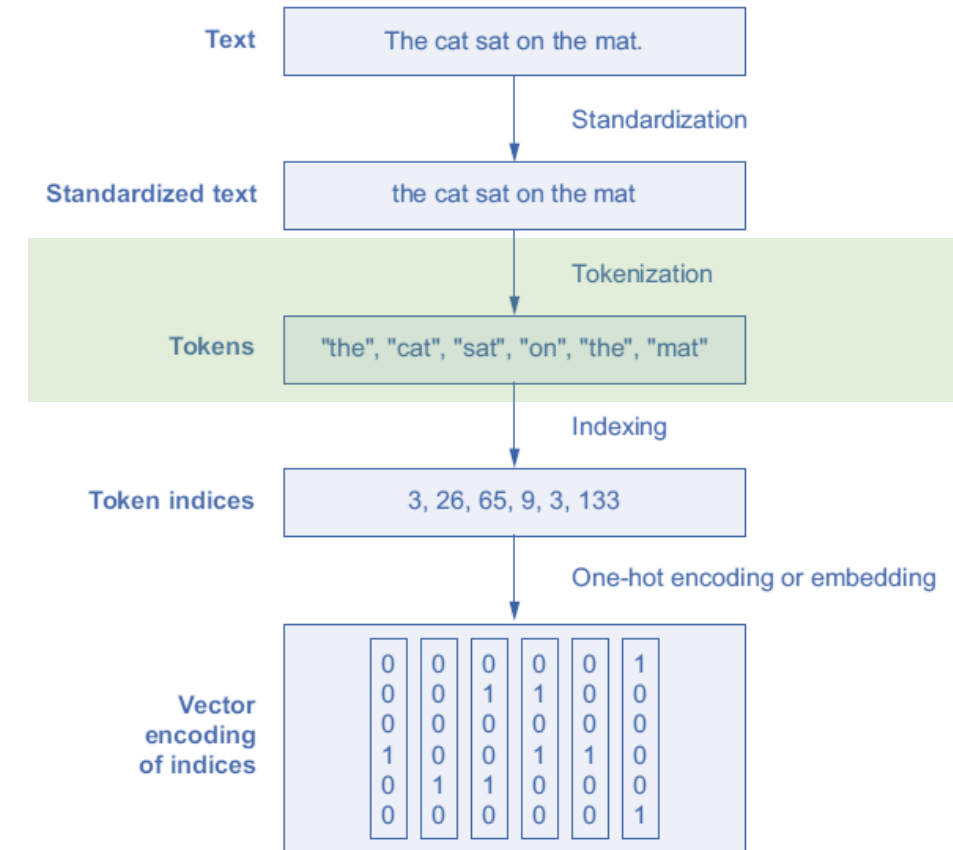
- **Standardization:**
  - Converting to *lowercase*,
  - Removing *punctuations*,
  - Converting *special characters*,
  - **Stemming** (converting variations of a term into a single shared representation)
- Standardization is a basic form of **feature engineering**
- With standardization, your model requires **less training data** and **generalize better**.





# Text Vectorization: standardization, **tokenization**, indexing, encoding

- **Tokenization: Splitting text into units/token**
  - *Word-level* tokenization: space-separated tokens
  - *N-gram* tokenization: groups of **N (or fewer)** consecutive words
  - *Character-level* tokenization: each character is a token
- Two kinds of text processing models:
  - **Bag-of-words** models (discarding original order)  
→ N-gram tokenization
  - **Sequence models** (word order matters)  
→ **word-level** tokenization

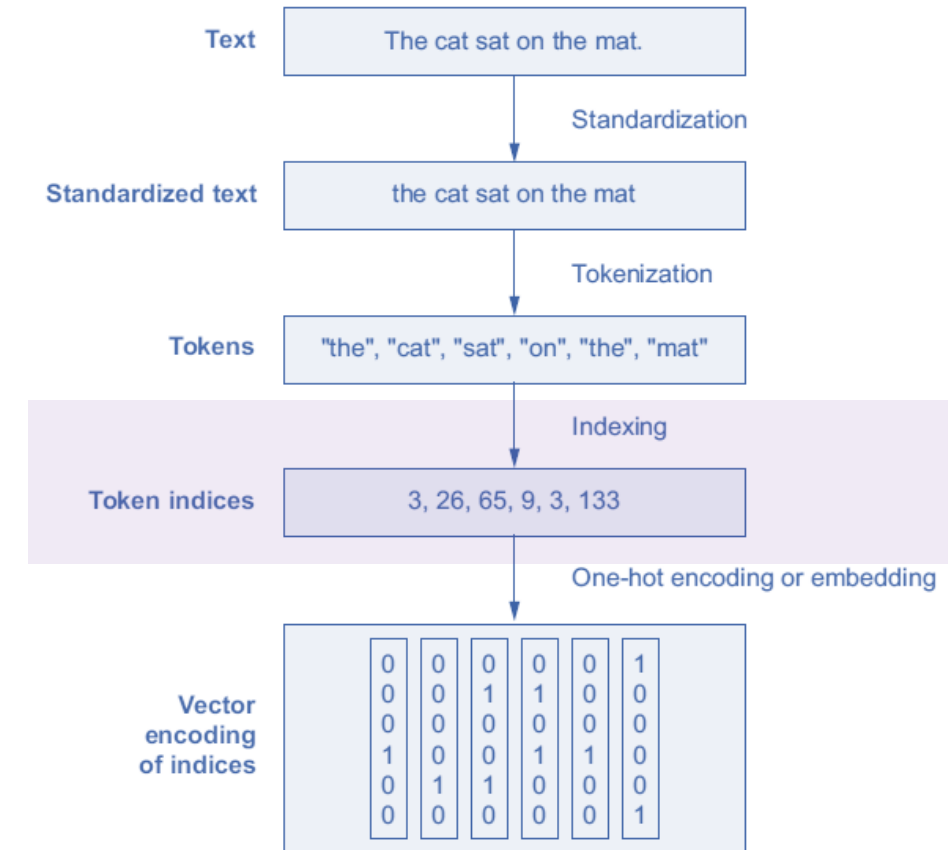




# Text Vectorization: standardization, tokenization, **indexing**, encoding

- **Indexing:**

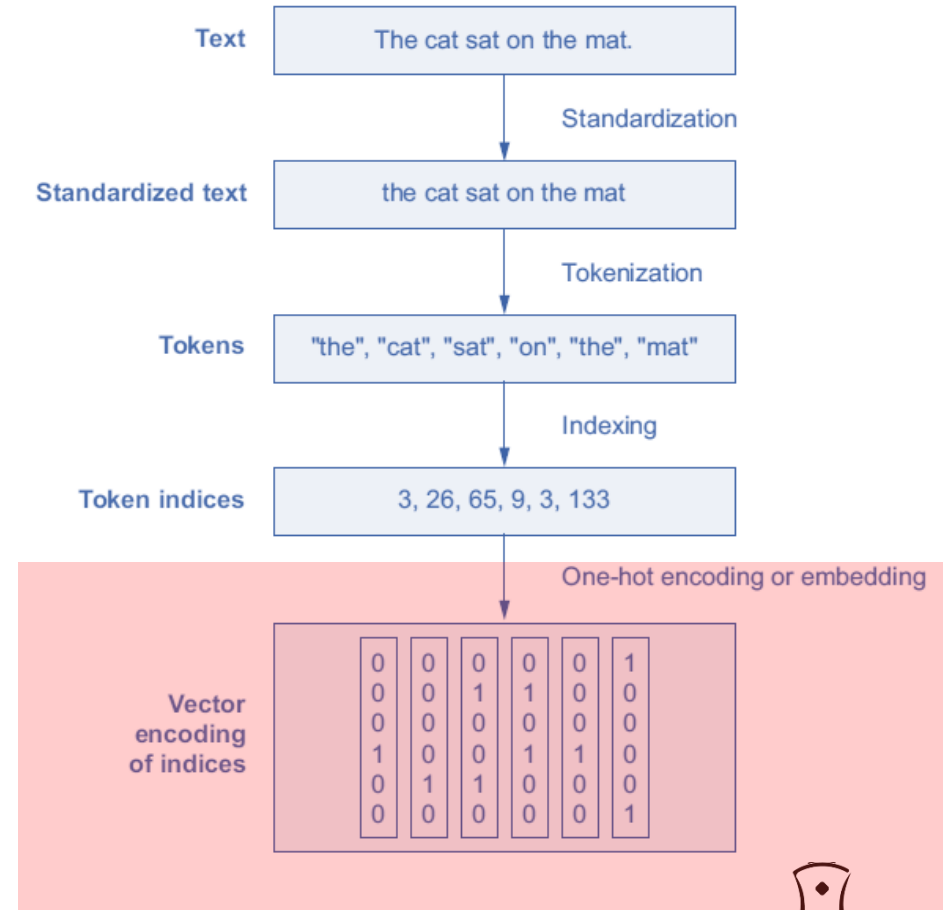
- **Integer**: return sequences of words encoded as integer indices
- Restrict to top 20k or 30k most common words
- **OOV** token (out of vocabulary) index **1**
- **Mask** token index **0** (ignore it, used for padding)



# Text Vectorization: standardization, tokenization, indexing, encoding

- **Encoding:**

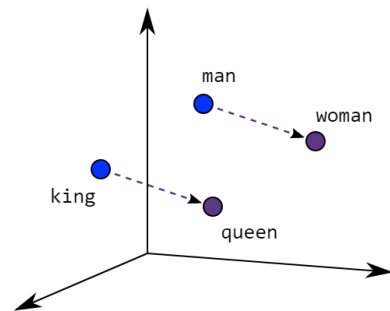
- **Multi-hot (binary):** Encode the tokens as a multi-hot binary vector
- **One-hot:** Encode the tokens as one-hot binary vectors
- **TF-IDF:** term frequency, inverse document frequency



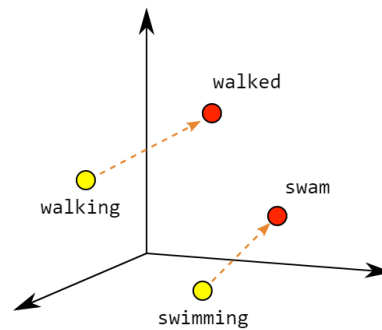


# Word Embedding

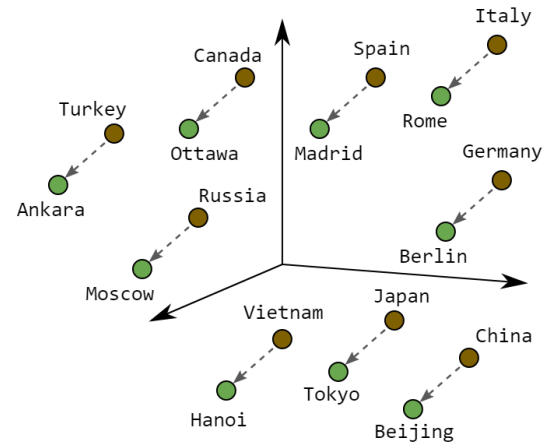
- Word embeddings are **dense**, **structured** representations **learned from data**.



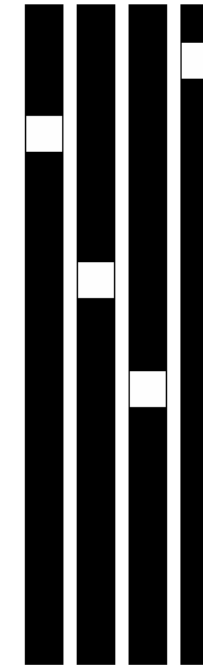
Male-Female



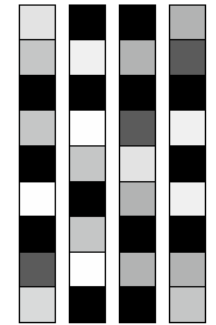
Verb Tense



Country-Capital



One-hot word vectors:  
- Sparse  
- High-dimensional  
- Hardcoded

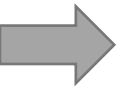


Word embeddings:  
- Dense  
- Lower-dimensional  
- Learned from data



# → Representing groups of words

	Bag-of-Words Approach	Sequential Modeling Approach
Input	Collection of text documents	Sequence of words
Output	Vector representation of each document	Sequential output, e.g. text classification, machine translation
Order	<b>Ignores</b> the order of words	<b>Captures</b> the order of words
Context	<b>Ignores</b> the context of words	<b>Captures</b> the context of words
Features	Each word is a feature	Features depend on the model architecture
Dimensionality	<b>High</b> dimensionality	<b>Lower</b> dimensionality
Sparsity	High sparsity	Lower sparsity
Performance	<b>Faster</b> training and inference	<b>Slower</b> training and inference
Suitability	Good for simple tasks like document classification	Good for complex tasks like language modeling and machine translation



# A simple example

- Historically, most early applications of machine learning to NLP involved bag-of-words models.
- Interest in sequence models started rising in 2015, with the rebirth of RNN
- Today, both approaches remain relevant. Let's look at a simple example: [IMDB movie review](#)
- Task: sentiment-classification of IMDB movie reviews (positive-negative)





## IMDB review example (bag-of-word approach)

- Single words (**Unigram**) with binary encoding (**multi-hot**)
- You can represent the entire text as a single vector!
- The cat sat on the mat  $\rightarrow$  {"cat", "mat", "sat", "on", "the"}  $\rightarrow$  (0,0,1,...,1,0,...,1,0,...,1,..1)
- Bag-of-words model is mainly made of **dense layers** (No recurrent layers)

```
from tensorflow import keras
from tensorflow.keras import layers

def get_model(max_tokens=20000, hidden_dim=16):
    inputs = keras.Input(shape=(max_tokens,))
    x = layers.Dense(hidden_dim, activation="relu")(inputs)
    x = layers.Dropout(0.5)(x)
    outputs = layers.Dense(1, activation="sigmoid")(x)
    model = keras.Model(inputs, outputs)
    model.compile(optimizer="rmsprop",
                  loss="binary_crossentropy",
                  metrics=["accuracy"])

    return model
```

```
text_vectorization = TextVectorization(
    max_tokens=20000,
    output_mode="multi_hot",
```

- Naïve baseline = **50%** (balance data)
- Test accuracy = **89.2%**



# IMDB review example (improving bag-of-word approach)

- We started with Single words (**Unigram**) with binary encoding (**multi-hot**)
- Two words (**Bigrams**) with binary encoding (**multi-hot**): adding **local order** information
- The cat sat on the mat  $\rightarrow$  {"the", "the cat", "cat", "cat sat", "sat", "sat on", "on", "on the", "the mat", "mat"}

## Test accuracy

- Naïve baseline = **50%** (balance data)
- Unigram binary = **89.2%**
- Bigram binary = **90.4%**

```
text_vectorization = TextVectorization(  
    ngrams=2,  
    max_tokens=20000,  
    output_mode="multi_hot",
```



# IMDB review example (improving bag-of-word approach)

- Bigrams with TF-IDF encoding

```
text_vectorization = TextVectorization(  
    ngrams=2,  
    max_tokens=20000,  
    output_mode="tf_idf",
```

Test accuracy

- Naïve baseline = 50% (balance data)
- Unigram binary = 89.2%
- Bigram binary = 90.4%
- Bigram TF-IDF = 89.8%
- Typically, TF-IDF increases the model performance by 1%. It wasn't the case for the IMDB dataset.



# IMDB review example (sequence modeling approach)

1. Starting by representing the input as sequences of **integer** indices (one integer for one word)
2. Map each integer to a vector to get **sequences of vectors** (one-hot or word embedding)
3. Finally, feed these vectors into a stack of layers such as **1D Convnet**, **RNN** or a **Transformer**.

```
from tensorflow.keras import layers

max_length = 600
max_tokens = 20000
text_vectorization = layers.TextVectorization(
    max_tokens=max_tokens,
    output_mode="int",
    output_sequence_length=max_length,
```



# IMDB review example (sequence modeling approach)

- One-hot encoding approach: Encode the integers into binary 20,000-dimensional vectors.

```
import tensorflow as tf
inputs = keras.Input(shape=(None,), dtype="int64")
embedded = tf.one_hot(inputs, depth=max_tokens)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.summary()
```

## Test accuracy

- Naïve baseline = 50% (balance data)
- Unigram binary = 89.2%
- Bigram binary = 90.4%
- Bigram TF-IDF = 89.8%
- Sequence one-hot encode = 87%

- This model trains very slowly.  $600 \times 20,000 = 12$  million floats for a single movie review!
- There must be a better way than one-hot encoding!





# IMDB review example (improving sequence modeling approach)

- Word-embedding approach:
  1. Simultaneously trained word embedding (trained jointly with the weights of a NN)
  2. Pretrained word embedding
- Sometimes it is reasonable to learn a new embedding space for different tasks (for example movie-review is different from legal-document classification)
- Word index → word embedding layer → corresponding word vector

```
inputs = keras.Input(shape=(None,), dtype="int64")
embedded = layers.Embedding(input_dim=max_tokens, output_dim=256)(inputs)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
model.compile(optimizer="rmsprop",
              loss="binary_crossentropy",
              metrics=["accuracy"])
```



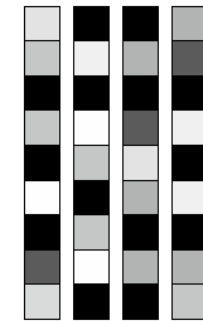


# IMDB review example (improving sequence modeling approach)

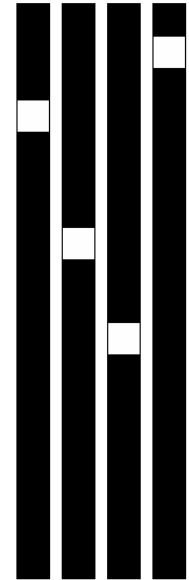
- **Word-embedding** approach performance:
- It trains much faster than the one-hot encoding (256 vs 20k dim)

## Test accuracy

- Naïve baseline = **50%** (balance data)
- Unigram binary = **89.2%**
- Bigram binary = **90.4%**
- Bigram TF-IDF = **89.8%**
- Sequence one-hot encode = **87%**
- Sequence word embedding simultaneously trained = **87%**
- Sequence word embedding simultaneously trained with **masks** = **88%**



Word embeddings:  
- Dense  
- Lower-dimensional  
- Learned from data



One-hot word vectors:  
- Sparse  
- High-dimensional  
- Hardcoded

# ➔ Pretrained word embedding

- Motivation: little training data! The network cannot learn an appropriate embedding.
- Analogous to using pretrained convnets in image classification
- Useful when expecting **generic** features.
- Schemes:
  - The **Word2Vec** algorithm (Tomas MiKolov at Google 2013): a **predictive model** that uses a neural network to predict the **probability of a word given its context**
  - **GloVe** (Stanford researchers 2014): a **count-based model** that computes the **co-occurrence probabilities** between words in a corpus
- Word2Vec is faster to train than GloVe, especially for large corpora

# ➔ Road map!

- ✓ Module 1- Introduction to Deep Forecasting
- ✓ Module 2- Setting up Deep Forecasting Environment
- ✓ Module 3- Exponential Smoothing
- ✓ Module 4- ARIMA models
- ✓ Module 5- Machine Learning for Time series Forecasting
- ✓ Module 6- Deep Neural Networks
- Module 7- Deep Sequence Modeling (RNN, LSTM)
- Module 8- Transformers (Attention is all you need!)
- Module 9- Prophet and Neural Prophet

